

CSC 261/461 – Database Systems

Lecture 20

Fall 2017

Announcements

- Project 3 (MongoDB) is out
 - Due on Dec 01
- Term paper is due on:
 - Dec 08, 2017
 - (You need to finish your poster before that to have ample time for getting it printed)
 - Details will follow...


Topics for Today

- MongoDB
- Query Processing (Chapter 18)
- Query Optimization (Chapter 19)

MONGODB

What is MongoDB

- Scalable High-Performance Open-source, Document-orientated database.
- Built for Speed
- Rich Document based queries for **Easy readability**.
- Full Index Support for **High Performance**.
- Map / Reduce for **Aggregation**.



Why use MongoDB?

- SQL was invented in the 70's to store **data**.
- MongoDB stores **documents (or) objects**
- Embedded documents and arrays reduce need for joins

Why will we use Mongodb?

- Semi-Structured Content Management

XML -> Tables

- Items -> User, Item, Category, Bid

Object-relational impedance mismatch

- A set of conceptual and technical difficulties that are often encountered:
 - when a **relational database management system (RDBMS)** is being served by an application program (or multiple application programs) written in an **object-oriented programming language**
- Objects or class definitions must be mapped to database tables defined by relational schema.

MongoDB: No Impedance Mismatch

```
// your application code
class Foo { int x; string [] tags; }

// mongo document for Foo
{ x: 1, tags: ['abc','xyz'] }
```

When I say
Database



Think
Database

- Made up of Multiple **Collections**.
- Created **on-the-fly** when referenced for the first time.

When I say
Collection



Think
Table

- Schema-less, and contains **Documents**.
- **Indexable** by one/more keys.
- Created **on-the-fly** when referenced for the first time.
- **Capped Collections:** Fixed size, older records get dropped after reaching the limit.

When I say
Document



Think
Record/Row

- Stored in a **Collection**.
- Have **_id** key – works like Primary keys in MySQL.
- Supported Relationships – **Embedded (or) References**.
- Document storage in **BSON** (Binary form of JSON).

The Document Model

```
var post = {  
    '_id': ObjectId('3432'),  
    'author': ObjectId('2311'),  
    'title': 'Introduction to MongoDB',  
    'body': 'MongoDB is an open sources.. ',  
    'timestamp': Date('01-04-12'),  
    'tags': ['MongoDB', 'NoSQL'],  
    'comments': [{  
        'author': ObjectId('5331'),  
        'date': Date('02-04-12'),  
        'text': 'Did you see.. ',  
        'upvotes': 7} ]  
}
```

```
> db.posts.insert(post);
```

Find

```
// find posts which has 'MongoDB' tag.
```

```
> db.posts.find({tags: 'MongoDB'});
```


```
// find posts by author's comments.
```

```
> db.posts.find({'comments.author': 'Johnson'}).count();
```

```
// find posts written after 31st March.
```

```
> db.posts.find({'timestamp': {'$gte': Date('31-03-12')} });
```

\$gt, \$lt, \$gte, \$lte, \$ne, \$all, \$in, \$nin...



Find

```
db.foo.find(query, projection)
```

Which fields?

Which documents?

Find: Projection

```
> db.posts.find({}, {title:1})
```

```
{ "_id" : ObjectId("5654381f37f63ffc4ebf1964"),  
  "title" : "NodeJS server" }  
{ "_id" : ObjectId("5654385c37f63ffc4ebf1965"),  
  "title" : "Introduction to MongoDB" }
```

Like

```
select title from posts
```

Empty projection like

```
select * from posts
```

Find

Find

- Query criteria
 - Single value field
 - Array field
 - Sub-document / dot notation

Projection

- Field inclusion and exclusion

Cursor

- Sort
- Limit
- Skip

Update

```
> db.posts.update(  
  {"_id" : ObjectId("5654381f37f63ffc4ebf1964")},  
  {  
    title:"NodeJS server"  
  });
```

This will **replace** the document by {title:"NodeJS server"}

Update: Change part of the document


```
> db.posts.update(  
  {"_id" : ObjectId("5654381f37f63ffc4ebf1964")},  
  {  
    $addToSet: {tags:"JS"},  
    $set: {title:"NodeJS server"},  
    $unset: { comments: 1}  
  });
```

\$set, \$unset

\$push, \$pull, \$pop, \$addToSet

\$inc, \$decr, many more...

Update



Options:


`{multi: true}` – will change all found documents;
by default only first found will be updated

`{upsert: true}` – will insert document if it was not found

Remove

- `db.collection.remove(<query>, <justOne>)`
- `db.items.remove({Currently: { $gt: 20 } })`

Aggregation





Aggregation

- <https://docs.mongodb.com/v3.0/applications/aggregation/>
- <https://www.safaribooksonline.com/blog/2013/06/21/aggregation-in-mongodb/>

MapReduce

Collection

```
db.orders.mapReduce(  
    map → function() { emit( this.cust_id, this.amount ); },  
    reduce → function(key, values) { return Array.sum( values ) },  
    query → { query: { status: "A" } },  
    output → { out: "order_totals" }  
)
```




query

map

{ "A123": [500, 250] } → reduce

{ "B212": 200 } →



order_totals

Acknowledgement

- Many of these slides are produced by Luxoft.com

QUERY PROCESSING

Steps in Query Processing

- Scanning
- Parsing
- Validation
- Query Tree Creation
- Query Optimization (Query planning)
- Code generation (to execute the plan)
- Running the query code

Nested Loop Joins

Notes

- We are again considering “IO aware” algorithms: *care about disk IO*
 - Given a relation R, let:
 - $T(R)$ = # of tuples in R
 - $P(R)$ = # of pages in R
 - Note also that we omit ceilings in calculations... good exercise to put back in!
- Recall that we read / write entire pages with disk IO

Nested Loop Join (NLJ)

Compute $R \bowtie S$ on A :

```
for r in R:  
    for s in S:  
        if r[A] == s[A]:  
            yield (r,s)
```

Nested Loop Join (NLJ)

Compute $R \bowtie S$ on A :

```
for r in R:  
    for s in S:  
        if r[A] == s[A]:  
            yield (r,s)
```

Cost:

$P(R)$

1. Loop over the tuples in R

Note that our IO cost is based on the number of ***pages*** loaded, not the number of tuples!

Nested Loop Join (NLJ)

Compute $R \bowtie S$ on A :

```
for r in R:  
    for s in S:  
        if r[A] == s[A]:  
            yield (r,s)
```

Cost:

$$P(R) + T(R)*P(S)$$

1. Loop over the tuples in R
2. For every tuple in R , loop over all the tuples in S

Have to read *all of S* from disk for *every tuple in R!*

Nested Loop Join (NLJ)

Compute $R \bowtie S$ on A :

```
for r in R:  
    for s in S:  
        if r[A] == s[A]:  
            yield (r,s)
```

Cost:

$$P(R) + T(R)*P(S)$$

1. Loop over the tuples in R
2. For every tuple in R , loop over all the tuples in S
3. **Check against join conditions**

Note that NLJ can handle things other than equality constraints... just check in the *if* statement!

Nested Loop Join (NLJ)

Compute $R \bowtie S$ on A :

```
for r in R:  
    for s in S:  
        if r[A] == s[A]:  
            yield (r,s)
```

Cost:

$$P(R) + T(R)*P(S) + OUT$$

1. Loop over the tuples in R
2. For every tuple in R , loop over all the tuples in S
3. Check against join conditions
4. Write out (to page, then when page full, to disk)

Nested Loop Join (NLJ)

Compute $R \bowtie S$ on A :

```
for r in R:  
    for s in S:  
        if r[A] == s[A]:  
            yield (r,s)
```

Cost:

$$P(R) + T(R)*P(S) + OUT$$

What if R ("outer") and S ("inner") switched?



$$P(S) + T(S)*P(R) + OUT$$

Outer vs. inner selection makes a huge difference-
DBMS needs to know which relation is smaller!

Block Nested Loop Join (BNLJ)

Block Nested Loop Join (BNLJ)

Compute $R \bowtie S$ on A :

```
for each page pr of R:  
    for page ps of S:  
        for each tuple r in pr:  
            for each tuple s in ps:  
                if r[A] == s[A]:  
                    yield (r,s)
```

Given **3** pages of memory

Cost:

$P(R)$

1. Load in 1 page of R at a time
(leaving 1 page each free for S & output)

Note: There could be some speedup here due to the fact that we're reading in multiple pages sequentially however we'll ignore this here!

Block Nested Loop Join (BNLJ)

Compute $R \bowtie S$ on A :

for each page pr of R :

 for page ps of S :

 for each tuple r in pr :

 for each tuple s in ps :

 if $r[A] == s[A]$:

 yield (r, s)

Given **3** pages of memory

Cost:

$$P(R) + P(R).P(S)$$

1. Load in 1 page of R at a time
(leaving 1 page each free for S & output)

2. **For each page segment of R ,
load each page of S**

Note: Faster to iterate over
the *smaller* relation first!

Block Nested Loop Join (BNLJ)

Compute $R \bowtie S$ on A :

```
for each page pr of R:  
    for page ps of S:  
        for each tuple r in pr:  
            for each tuple s in ps:  
                if r[A] == s[A]:  
                    yield (r,s)
```

Given **3** pages of memory

Cost:

$$P(R) + P(R).P(S)$$

1. Load in 1 page of R at a time (leaving 1 page each free for S & output)
2. For each page segment of R, load each page of S
3. Check against the join conditions

BNLJ can also handle non-equality constraints

Block Nested Loop Join (BNLJ)

Compute $R \bowtie S$ on A :

```
for each page pr of R:  
    for page ps of S:  
        for each tuple r in pr:  
            for each tuple s in ps:  
                if r[A] == s[A]:  
                    yield (r,s)
```

Given 3 pages of memory

Cost:

$$P(R) + P(R).P(S)$$

1. Load 1 page of R at a time
(leaving 1 page each free for S & output)
2. For each page segment of R,
load each page of S
3. Check against the join
conditions
4. Write out

Block Nested Loop Join (BNLJ) (B+1 pages of Memory)

Compute $R \bowtie S$ on A :

```
for each B-1 pages pr of R:  
    for page ps of S:  
        for each tuple r in pr:  
            for each tuple s in ps:  
                if r[A] == s[A]:  
                    yield (r,s)
```

Given **B+1** pages of memory

Cost:

$P(R)$

1. Load in $B-1$ pages of R at a time (leaving 1 page each free for S & output)

Note: There could be some speedup here due to the fact that we're reading in multiple pages sequentially however we'll ignore this here!

Block Nested Loop Join (BNLJ)

Compute $R \bowtie S$ on A :

for each $B-1$ pages pr of R :

 for page ps of S :

 for each tuple r in pr :

 for each tuple s in

ps :

 if $r[A] == s[A]$:
 yield (r, s)

Given $B+1$ pages of memory

Cost:

$$P(R) + \frac{P(R)}{B-1} P(S)$$

1. Load in $B-1$ pages of R at a time (leaving 1 page each free for S & output)
2. For each $(B-1)$ -page segment of R , load each page of S

Note: Faster to iterate over the *smaller* relation first!

Block Nested Loop Join (BNLJ)

Compute $R \bowtie S$ on A :

```
for each  $B-1$  pages  $pr$  of  $R$ :  
    for page  $ps$  of  $S$ :  
        for each tuple  $r$  in  $pr$ :  
            for each tuple  $s$  in  
                 $ps$ :
```

```
        if  $r[A] == s[A]$ :  
            yield ( $r, s$ )
```

Given $B+1$ pages of memory

Cost:

$$P(R) + \frac{P(R)}{B-1} P(S)$$

1. Load in $B-1$ pages of R at a time (leaving 1 page each free for S & output)
2. For each $(B-1)$ -page segment of R , load each page of S
3. Check against the join conditions

BNLJ can also handle non-equality constraints

Block Nested Loop Join (BNLJ)

Compute $R \bowtie S$ on A :

```
for each  $B-1$  pages pr of R:  
    for page ps of S:  
        for each tuple r in pr:  
            for each tuple s in  
                ps:  
                    if  $r[A] == s[A]$ :  
                        yield (r, s)
```

Given $B+1$ pages of memory

Cost:

$$P(R) + \frac{P(R)}{B-1} P(S) + \text{OUT}$$

1. Load in $B-1$ pages of R at a time (leaving 1 page each free for S & output)
2. For each $(B-1)$ -page segment of R, load each page of S
3. Check against the join conditions
4. Write out

BNLJ vs. NLJ: Benefits of IO Aware

- In BNLJ, by loading larger chunks of R, we minimize the number of full *disk reads* of S
 - We only read all of S from disk for *every $(B-1)$ -page segment of R!*
 - Still the full cross-product, but more done only *in memory*

NLJ

$$P(R) + T(R)*P(S) + OUT$$



BNLJ

$$P(R) + \frac{P(R)}{B-1} P(S) + OUT$$

BNLJ is faster by roughly $\frac{(B-1)T(R)}{P(R)}$

BNLJ vs. NLJ: Benefits of IO Aware

- Example:
 - R: 500 pages
 - S: 1000 pages
 - 100 tuples / page
 - We have 12 pages of memory ($B = 11$)
- NLJ: Cost = $500 + 50,000 * 1000 = 50 \text{ Million IOs} \sim= \underline{140 \text{ hours}}$
- BNLJ: Cost = $500 + \frac{500 * 1000}{10} = 50 \text{ Thousand IOs} \sim= \underline{0.14 \text{ hours}}$

Ignoring OUT here...

A very real difference from a small
change in the algorithm!

Smarter than Cross-Products

Smarter than Cross-Products: From Quadratic to Nearly Linear

- All joins that compute the *full cross-product* have some **quadratic** term
 - For example we saw: $\text{NLJ } P(R) + \textcolor{red}{T(R)P(S)} + OUT$
 - $\text{BNLJ } P(R) + \frac{\textcolor{red}{P(R)}}{B-1} \textcolor{red}{P(S)} + OUT$
- Now we'll see some (nearly) linear joins:
 - $\sim O(P(R) + P(S) + OUT)$, where again *OUT* could be quadratic but is usually better

We get this gain by *taking advantage of structure*- moving to equality constraints (“equijoin”) only!

Index Nested Loop Join (INLJ)

Compute $R \bowtie S$ on A :

Given index idx on
 $S.A$:

```
for r in R:  
    s in idx(r[A]):  
        yield r, s
```

Cost:

$$P(R) + T(R)*L + OUT$$

where L is the IO cost to access all the distinct values in the index; assuming these fit on one page, $L \sim 3$ is good est.

→ We can use an **index** (e.g. B+ Tree) to ***avoid doing the full cross-product!***

Sort-Merge Join (SMJ)

What you will learn about in this section

1. Sort-Merge Join
2. “Backup” & Total Cost
3. Optimizations

Sort Merge Join (SMJ): Basic Procedure

To compute $R \bowtie S$ on A :


1. Sort R, S on A using *external merge sort*
2. *Scan* sorted files and “merge”
3. [May need to “backup” - see next subsection]

Note that we are only considering equality join conditions here

Note that if R, S are already sorted on A ,
SMJ will be awesome!


SMJ Example: $R \bowtie S$ on A with 3 page buffer

- For simplicity: Let each page be *one tuple*, and let the first value be A




SMJ Example: $R \bowtie S$ on A with 3 page buffer

- i. Sort the relations R, S on the join key (first value)




SMJ Example: $R \bowtie S$ on A with 3 page buffer

2. Scan and “merge” on join key!




SMJ Example: $R \bowtie S$ on A with 3 page buffer

2. Scan and “merge” on join key!




SMJ Example: $R \bowtie S$ on A with 3 page buffer

2. Scan and “merge” on join key!



SMJ Example: $R \bowtie S$ on A with 3 page buffer


2. Done!



What happens with duplicate join keys?


Multiple tuples with Same Join Key: “Backup”

i. Start with sorted relations, and begin scan / merge...




Multiple tuples with Same Join Key: “Backup”

i. Start with sorted relations, and begin scan / merge...




Multiple tuples with Same Join Key: “Backup”

i. Start with sorted relations, and begin scan / merge...



Multiple tuples with Same Join Key: “Backup”

i. Start with sorted relations, and begin scan / merge...



Backup

- At best, no backup \rightarrow scan takes $P(R) + P(S)$ reads
 - For ex: if no duplicate values in join attribute
- At worst (e.g. full backup each time), scan could take $P(R) * P(S)$ reads!
 - For ex: if *all* duplicate values in join attribute, i.e. all tuples in R and S have the same value for the join attribute
 - Roughly: For each page of R, we'll have to *back up* and read each page of S...
- Often not that bad however

SMJ: Total cost

- Cost of SMJ is **cost of sorting** R and S...
- Plus the **cost of scanning**: $\sim P(R) + P(S)$
 - Because of *backup*: in worst case $P(R)*P(S)$; but this would be very unlikely
- Plus the **cost of writing out**: $\sim P(R) + P(S)$ but in worst case $T(R)*T(S)$

$\sim \text{Sort}(P(R)) + \text{Sort}(P(S))$
 $+ P(R) + P(S) + \text{OUT}$


SMJ vs. BNLJ

- If we have **100** buffer pages, $P(R) = 1000$ pages and $P(S) = 500$ pages:
 - Sort both in two passes: $2 * 2 * 1000 + 2 * 2 * 500 = 6,000$ IOs
 - Merge phase $1000 + 500 = 1,500$ IOs
 - **7,500 IOs + OUT**

What is BNLJ?

- $500 + 1000 * \left\lceil \frac{500}{98} \right\rceil = \underline{\b{6,500 IOs + OUT}}$
- But, if we have **35** buffer pages?
 - Sort Merge has same behavior (still 2 passes)
 - BNLJ? **15,500 IOs + OUT!**

Basic SMJ



Takeaway points from SMJ

If input already sorted on join key, skip the sorts.

- SMJ is basically linear.
- Nasty but unlikely case: Many duplicate join keys.

4. HASH JOIN (HJ)

What you will learn about in this section

1. Hash Join
2. Memory requirements

Recall: Hashing

- **Magic of hashing:**
 - A hash function h_B maps into $[0, B-1]$
 - And maps nearly uniformly
- A hash **collision** is when $x \neq y$ but $h_B(x) = h_B(y)$
 - Note however that it will never occur that $x = y$ but $h_B(x) \neq h_B(y)$

Hash Join: High-level procedure

To compute $R \bowtie S$ on A :


Note again that we are only considering equality constraints here

1. **Partition Phase:** Using one (shared) hash function h_B , partition R and S into B buckets
2. **Matching Phase:** Take pairs of buckets whose tuples have the same values for h , and join these
 - i. Use BNLJ here; or hash again → either way, operating on small partitions so fast!

We **decompose** the problem using h_B , then complete the join

Hash Join: High-level procedure

1. Partition Phase: Using one (shared) hash function h_B , partition R and S into B buckets





Note our new convention:
pages each have two tuples (one per row)

More detail in a second...


Hash Join: High-level procedure

2. Matching Phase: Take pairs of buckets whose tuples have the same values for h_B , and join these



Hash Join: High-level procedure

2. Matching Phase: Take pairs of buckets whose tuples have the same values for h_B , and join these



Hash Join Phase 1: Partitioning

Goal: For each relation, partition relation into **buckets** such that if $h_B(t_i.A) = h_B(t_j.A)$ they are in the same bucket

Given $B+i$ buffer pages, we partition into B buckets:

- We use B buffer pages for output (one for each bucket), and i for input
 - For each tuple t in input, copy to buffer page for $h_B(t.A)$
 - When page fills up, flush to disk.

How big are the resulting buckets?

Given $B+1$ buffer
pages

- Given N input pages, we partition into B buckets:
 - → Ideally our buckets are each of size $\sim N/B$ pages

How big *do we want* the resulting buckets?

- Ideally, our buckets would be of size $\leq B - 1$ **pages**
 - **1** for input page, **1** for output page, **$B-1$** for each bucket
- Recall: If we want to join a bucket from R and one from S, we can do BNLJ **in linear time** if for *one of them* (*wlog say R*), $P(R) \leq B - 1$!
 - And more generally, being able to fit bucket in memory is advantageous
- We can keep partitioning buckets that are $> B-1$ pages, until they are $\leq B - 1$ **pages**
 - Using a new hash key which will split them...

Given **$B+1$** buffer pages


Recall for BNLJ:
 $P(R)$
 $+ \frac{P(R)P(S)}{B - 1}$

We'll call each of these a "pass" again...

Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages

We partition into $B = 2$ buckets using hash function h_2 so that we can have one buffer page for each partition (and one for input)




For simplicity, we'll look at partitioning one of the two relations- we just do the same for the other relation!

Recall: our goal will be to get $B = 2$ buckets of size $\leq B-1 \rightarrow 1$ page each

Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages


- We read pages from R into the “input” page of the buffer...



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages


2. Then we use hash function h_2 to sort into the buckets, which each have one page in the buffer



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages


2. Then we use **hash function h_2** to sort into the buckets, which each have one page in the buffer



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages


3. We repeat until the buffer bucket pages are full...



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages


3. We repeat until the buffer bucket pages are full...



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages


3. We repeat until the buffer bucket pages are full...



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages


3. We repeat until the buffer bucket pages are full...
then flush to disk



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages



3. We repeat until the buffer bucket pages are full...
then flush to disk



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages



Note that collisions can occur!



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages


Finish this pass...



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages



Finish this pass...



Hash Join Phase 1: Partitioning

Finish this pass...


Given $B+1 = 3$ buffer pages



Hash Join Phase 1: Partitioning


Given $B+1 = 3$ buffer pages

Finish this pass...



Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages




We wanted buckets of size $B-1 = 1$... *however we got larger ones due to:*

- (1) Duplicate join keys
- (2) Hash collisions

Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages




To take care of larger buckets caused by (2) hash collisions, we can just do another pass!
What hash function should we use?

Do another pass with a different hash function, h'_2 , ideally such that:

$$h'_2(3) \neq h'_2(5)$$

Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages




To take care of larger buckets caused by (2) hash collisions, we can just do another pass!
What hash function should we use?

Do another pass with a different hash function, h'_2 , ideally such that:

$$h'_2(3) \neq h'_2(5)$$

Hash Join Phase 1: Partitioning

Given $B+1 = 3$ buffer pages




What about duplicate join keys?
Unfortunately this is a problem...
but usually not a huge one.

We call this unevenness
in the bucket size skew

Now that we have partitioned R and S...

Hash Join Phase 2: Matching

- Now, we just join pairs of buckets from R and S that have the same hash value to complete the join!




Hash Join Phase 2: Matching

- Note that since $x = y \rightarrow h(x) = h(y)$, we only need to consider pairs of buckets (one from R, one from S) that have the same hash function value
- If our buckets are $\sim B - 1$ pages, can join each such pair using BNLJ *in linear time*; recall (with $P(R) = B-1$):


$$\text{BNLJ Cost: } P(R) + \frac{P(R)P(S)}{B-1} = P(R) + \frac{(B-1)P(S)}{B-1} = P(R) + P(S)$$

Joining the pairs of buckets is linear!
(As long as smaller bucket $\leq B-1$ pages)

Hash Join Phase 2: Matching



Hash Join Phase 2: Matching




$R \bowtie S \text{ on } A$

To perform the join, we ideally just need to explore the dark blue regions

= the tuples with same values of the join key A


Hash Join Phase 2: Matching



$R \bowtie S \text{ on } A$

With a join algorithm like BNLJ that doesn't take advantage of equijoin structure, we'd have to explore this **whole grid!**

Hash Join Phase 2: Matching



$R \bowtie S \text{ on } A$


With HJ, we only explore the **blue** regions

= the tuples with same values of $h(A)$!

We can apply BNLJ to each of these regions

Hash Join Phase 2: Matching

R.A
hashed
values



$R \bowtie S \text{ on } A$

An alternative to
applying BNLJ:

We could also hash
again, and keep
doing passes in
memory to reduce
further!

Hash Join Summary

- **Partitioning** requires reading + writing each page of R,S
 - $\rightarrow 2(P(R)+P(S))$ IOs
- **Matching** (with BNLJ) requires reading each page of R,S
 - $\rightarrow P(R) + P(S)$ IOs
- **Writing out results** could be as bad as $P(R)*P(S)...$ but probably closer to $P(R)+P(S)$

HJ takes $\sim 3(P(R)+P(S)) + OUT$ IOs!

Sort-Merge vs. Hash Join

- *Given enough memory*, both SMJ and HJ have performance:

$$\sim 3(P(R) + P(S)) + \\ OUT$$

- “Enough” memory =

- SMJ: $B^2 > \max\{P(R), P(S)\}$

- HJ: $B^2 > \min\{P(R), P(S)\}$

Hash Join superior if relation sizes ***differ greatly***. Why?

Further Comparisons of Hash and Sort Joins

- Hash Joins are highly parallelizable.
- Sort-Merge less sensitive to data skew and result is sorted

Summary

- Saw IO-aware join algorithms
 - Massive differences in performance.

Acknowledgement

- Some of the slides in this presentation are taken from the slides provided by the authors.
- Many of these slides are taken from cs145 course offered by Stanford University.